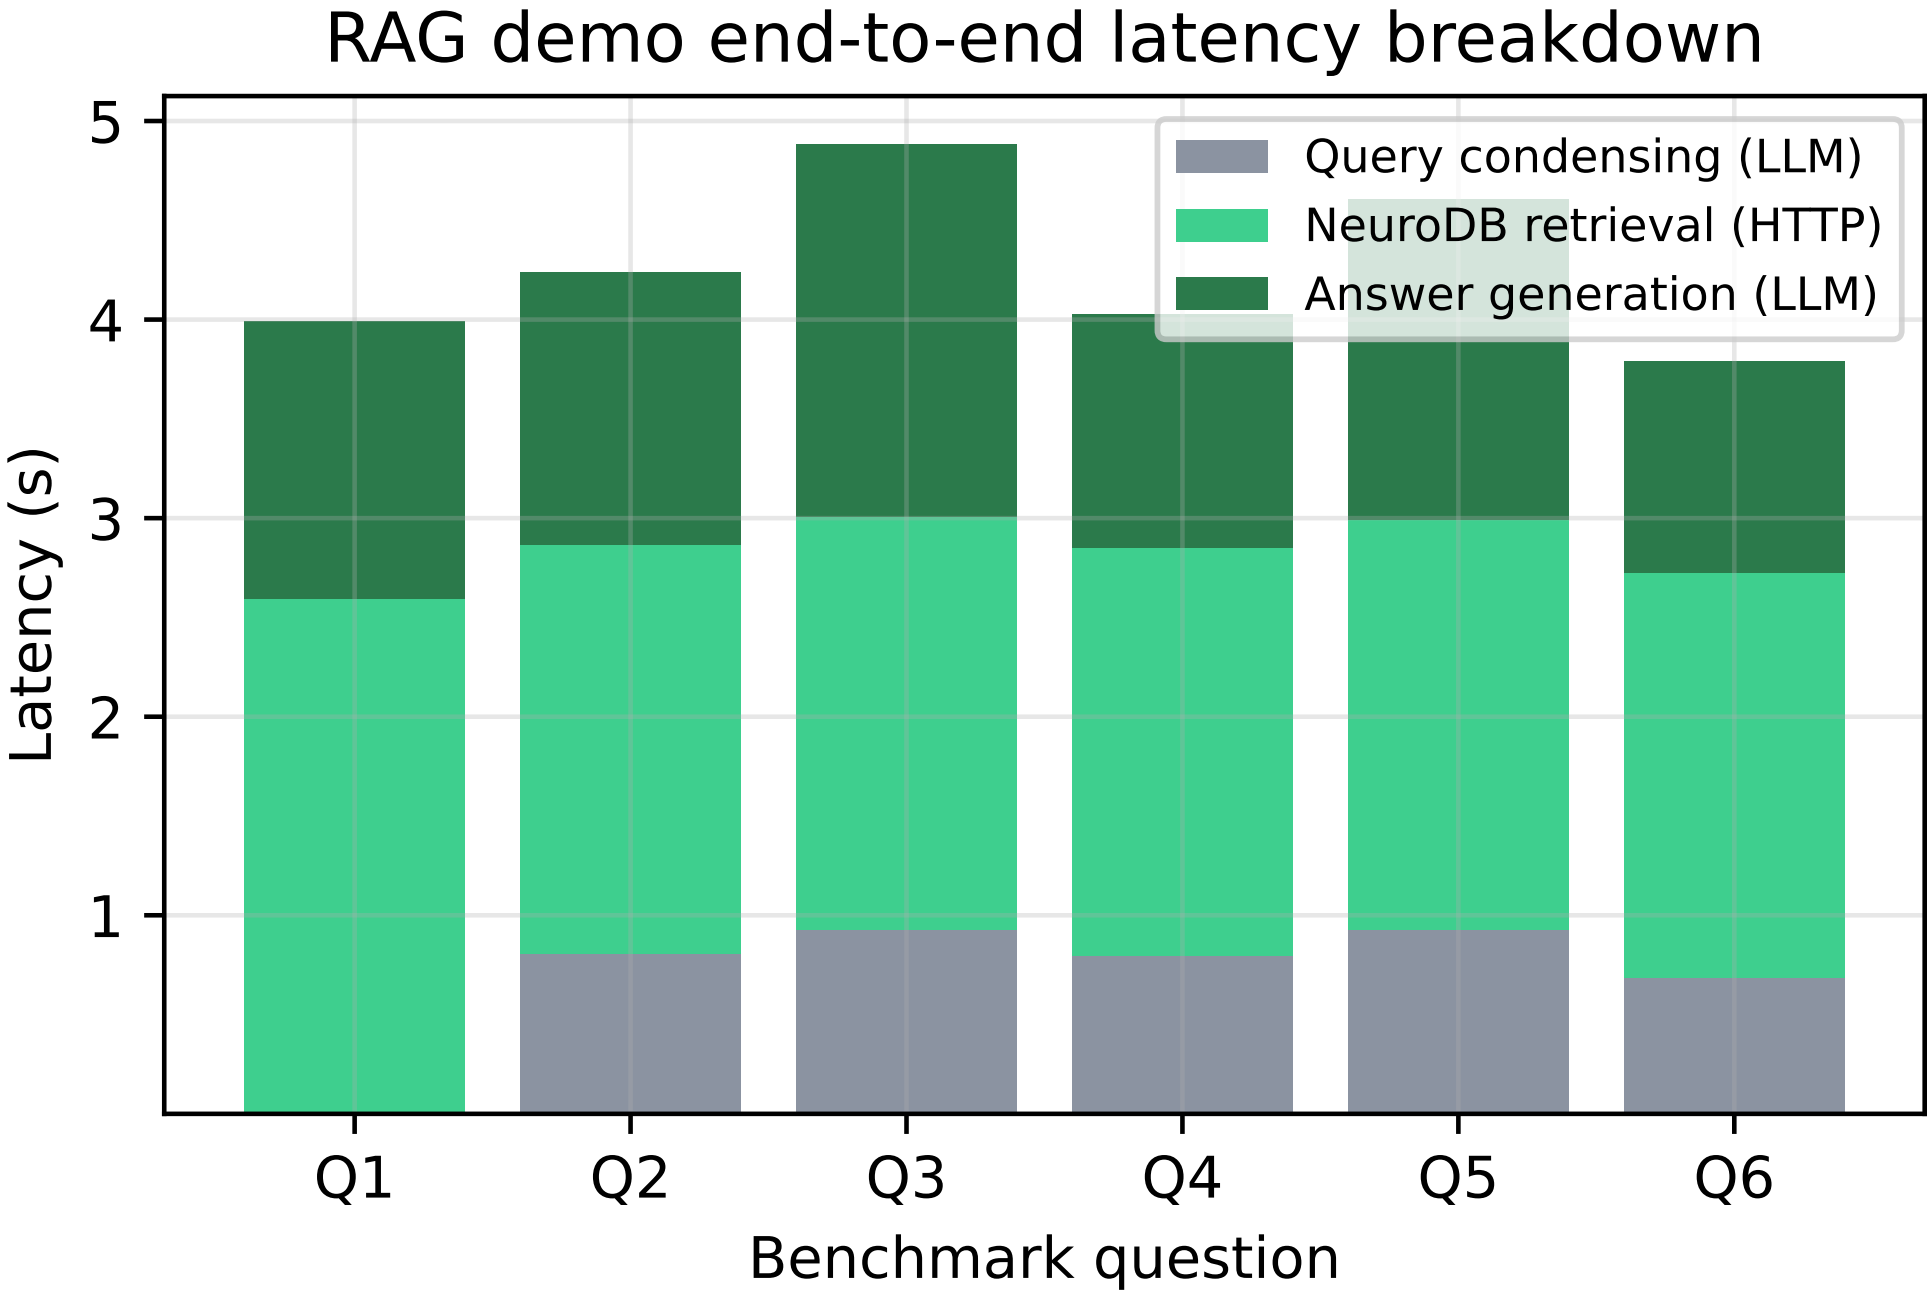
RAG demo end-to-end latency breakdown
Query condensing (LLM)
NeuroDB retrieval (HTTP)
Answer generation (LLM)
Latency (s)
5
4
3
2
1
Q1
Q2
Q3
Q4
Q5
Q6
Benchmark question